

Fabian Bicca Moraes

507 778 854 | fabianmoraesb@gmail.com | <https://www.linkedin.com/in/fabian-bicca-59248b21b/> | github.com/Biccamor

EDUCATION

AGH University of Science and Technology

Kraków, Poland

B.Sc. Eng. in Computer Science (Inżynier) – Faculty of IT (WI)

Oct. 2025 – Present

- Relevant Coursework: Linear Algebra, Calculus, Algorithms & Data Structures, Discrete Maths.

EXPERIENCE & ACTIVITIES

BIT Student Computer Science Association

Kraków, Poland

Active Member, AI Section

March 2026 – Present

- Developing an **Intelligent CV Parsing Pipeline** using **Qwen 3.6** and **Pydantic** for structured data extraction.
- Presented research at the **62nd Student Scientific Conference (AGH)**, contributing to a **3rd place** award-winning technical project in applied computer science section.

EnsembleAI Hackathon

Kraków, Poland

Infrastructure & AI Task Architect

Jan. 2026 – March 2026

- Designed **HPC-optimized** infrastructure for LLM tasks, managing resources for 200 participants.
- Automated evaluation scripts to ensure stable inter-module communication and real-time scoring.

PROJECTS

Moovly – Group AI Recommendations | *FastAPI, RAG, pgvector, Llama 3.1*

Feb 2026 – Present

- Built a **RAG pipeline** using **BGE-en-v1.5** and **FlashRank** reranking for group-based movie suggestions.
- Developed a custom **weighted embedding logic** based on a predefined mapping dictionary to align user "vibes" with specific genres and keywords, improving recommendation relevance.
- Integrated **pgvector/PostgreSQL** with **SlowAPI** rate limiting to manage **Groq API** quotas.
- Developed the robust AI backend architecture and successfully integrated it with a **React frontend** to deliver a seamless user experience.
- Containerized the application using **Docker** and currently establishing a Continuous Deployment (CD) workflow targeting Google Cloud Platform (**GCP**) for automated, cloud-native hosting.
- Designed and implemented a **CI/CD** pipeline using GitHub Actions, automating code quality checks (linter), security scanning (Bandit), and regression testing (pytest) with build caching for optimized execution times.

Multimodal Video Search Engine | *SigLIP, Faster-Whisper, Streamlit, MiniLM*

Dec. 2025 – Feb. 2026

- Developed a tool for semantic video retrieval by aligning text queries with visual and audio embeddings.
- Optimized performance via **frame-sampling** and **model quantization** for efficient real-time inference.

Intelligent CV Parsing Pipeline | *Python, Qwen 3.6, GLM-OCR, PyMuPDF4LLM*

March 2026 – Present

- Engineered a secure extraction layer using Qwen to generate PII-free JSONs with **Pydantic**, ensuring safe and anonymized data hand-off to the team's downstream feedback agents.
- Project serves as a core engine for a BIT science club's automated student feedback system.
- Containerized the extraction pipeline for scalable deployment and robust inter-module communication.

SKILLS & LANGUAGES

Tools & Backend: Python, FastAPI, Docker, PostgreSQL, Pytest, Pydantic, Git, HPC, SlowAPI, Github Actions.

AI/ML: RAG, LLMs, pgvector, LanceDB, Embeddings (SigLIP, BGE), FlashRank, PyTorch, S-BERT.

Data Science: Pandas, Numpy.

Languages: Polish (Native), English (**IELTS C1 Academic**), Spanish (A2).

Wyrażam zgodę na przetwarzanie moich danych osobowych dla potrzeb niezbędnych do realizacji procesu rekrutacji (zgodnie z ustawą z dnia 10 maja 2018 roku o ochronie danych osobowych (Dz. U. z 2018, poz. 1000) oraz zgodnie z Rozporządzeniem Parlamentu Europejskiego i Rady (UE) 2016/679 z dnia 27 kwietnia 2016 r. (RODO)).